

Terraform: Cloud Configuration Management


Martin Schütte

18 April 2017




Concepts

From Servers ...



by Rodzilla at [Wikimedia Commons](#) (CC-BY-SA-3.0)

...to Services



Services also need Configuration Management

- Replace “click paths” with source code in VCS
 - Lifecycle awareness, not just a `setup.sh`
 - Reproducible environments
 - Specification, documentation, policy enforcement
- ⇒ Infrastructure as Code



TERRAFORM

Build, Combine, and Launch Infrastructure



Example: Simple Webservice (part 1)

```
### AWS Setup
provider "aws" {
  profile = "${var.aws_profile}"
  region  = "${var.aws_region}"
}

# Queue
resource "aws_sqs_queue" "importqueue" {
  name = "${var.app_name}-${var.aws_region}-importqueue"
}

# Storage
resource "aws_s3_bucket" "importdisk" {
  bucket = "${var.app_name}-${var.aws_region}-importdisk"
  acl     = "private"
}
```

Example: Simple Webservice (part 2)

```
### Heroku Setup
provider "heroku" { ... }

# Importer
resource "heroku_app" "importer" {
  name      = "${var.app_name}-${var.aws_region}-import"
  region    = "eu"
  config_vars {
    SQS_QUEUE_URL = "${aws_sqs_queue.importqueue.id}"
    S3_BUCKET      = "${aws_s3_bucket.importdisk.id}"
  }
}

resource "heroku_addon" "mongolab" {
  app      = "${heroku_app.importer.name}"
  plan     = "mongolab:sandbox"
}
```


- Simple model of resource entities with attributes
- Stateful lifecycle with CRUD operations
- Declarative configuration
- Dependencies by inference
- Parallel execution

- Provider: a source of resources
(usually with an API endpoint & authentication)
- Resource: every thing “that has a set of configurable attributes and a lifecycle (create, read, update, delete)” – implies ID and state
- Data Source: information read from provider
(e.g. lookup own account ID or AMI-ID)
- Provisioner: initialize a resource with local or remote scripts

- Order: directed acyclic graph of all resources
- Plan: generate an execution plan for review before applying a configuration
- State: execution result is kept in state file (local or remote)
- Lightweight: little provider knowledge, no error handling

Providers:

- **AWS**
- **Azure**
- **Google Cloud**
- Heroku
- DNSMadeEasy
- OpenStack
- Docker
- ...

Resources:

- azurerm_lb
- azurerm_subnet
- azurerm_dns_zone
- azure_instance
- aws_iam_user
- heroku_app
- postgresql_schema
- ...

Provisioners:

- chef
- file
- local-exec
- remote-exec

- Hashicorp Configuration Language (HCL), think “JSON-like but human-friendly”
- Variables
- Interpolation, e.g.
`"number ${count.index + 1}"`
- Attribute access with `resource_type.resource_name`
- Few build-in functions, e.g.
`base64encode(string), format(format, args...)`

HCL vs. JSON

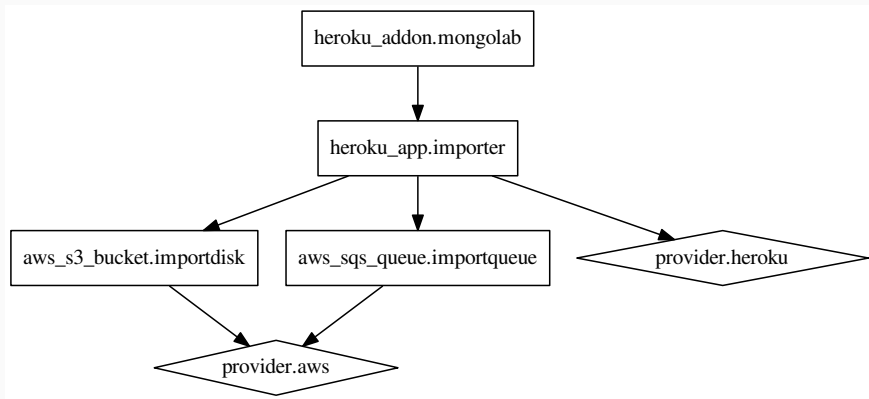
```
# An AMI
variable "ami" {
    description = "custom AMI"
}

/* A multi
   line comment. */
resource "aws_instance" "web" {
    ami = "${var.ami}"
    count = 2
    source_dest_check = false

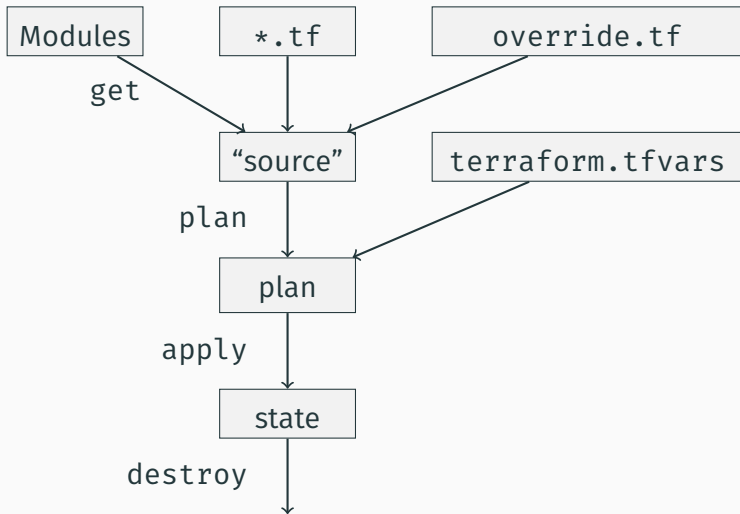
    connection {
        user = "root"
    }
}
```

```
{
  "variable": {
    "ami": {
      "description": "custom AMI"
    }
  },
  "resource": {
    "aws_instance": {
      "web": {
        "ami": "${var.ami}",
        "count": 2,
        "source_dest_check": false,

        "connection": {
          "user": "root"
        }
      }
    }
  }
}
```



Terraform Process



Example: Add Provisioning

```
# Importer
resource "heroku_app" "importer" {
  name      = "${var.app_name}-${var.aws_region}-import"
  region    = "eu"

  config_vars { ... }

  provisioner "local-exec" {
    command = <<EOT
cd ~/projects/go-testserver &&
git remote add heroku ${heroku_app.importer.git_url} &&
git push heroku master
EOT
  }
}
```

Example: Add Outputs

```
# Storage
resource "aws_s3_bucket" "importdisk" { ... }

# Importer
resource "heroku_app" "importer" { ... }

# Outputs
output "importer_bucket_arn" {
  value = "${aws_s3_bucket.importdisk.arn}"
}

output "importer_url" {
  value = "${heroku_app.importer.web_url}"
}

output "importer_gitrepo" {
  value = "${heroku_app.importer.git_url}"
}
```

Example: Add Lifecycle Meta-Parameter

```
# Storage
resource "aws_s3_bucket" "importdisk" {
  bucket = "${var.app_name}-${var.aws_region}-importdisk"
  acl     = "private"

  lifecycle {
    prevent_destroy = true
  }
}
```

```
$ terraform validate
$ terraform plan -out=my.plan
$ terraform show my.plan
$ terraform apply my.plan

$ terraform output
$ terraform output -json
$ terraform output importer_url
$ curl -s $(terraform output importer_url)

$ terraform graph | dot -Tpdf > graph.pdf && evince graph.pdf

$ terraform plan -destroy
$ terraform destroy
```

Features

“Plain terraform code” lacks structure and reusability

Modules

- are subdirectories with self-contained terraform code
- may be sourced from Git, Mercurial, HTTPS locations
- use variables and outputs to pass data

Example Module

```
module "vm" {  
  source =  
    "github.com/chiradeep/terraform-azure-modules/ubuntu_public_vm"  
  
  resource_group_name = "alpha-4ca23e"  
  location             = "West US"  
  name                 = "alpha"  
  
  vhd_uri_base        =  
    "https://alphastacctf62da8.blob.core.windows.net/alphacntrn22/"  
  user_data_file       = "userdata.sh"  
  subnet_id            = "subnet-public-98a1e34"  
  ssh_public_keyfile   = "key.pub"  
  
  vm_count = 2  
}
```

using [chiradeep/terraform-azure-modules](#) 

- Terraform keeps known state of resources
- Defaults to local state in `terraform.tfstate`
- Optional remote state with different backends (Azure Storage, S3, Consul, Atlas, ...)
 - Useful to sync multiple team members
 - May need additional mutex mechanism (v0.9 added state locking for Local, S3, and Consul)
 - Remote state is a data source

Example: Using State Import

```
$ terraform import azurerm_storage_account.my_storage_account \  
  /subscriptions/e9b2ec19-ab6e-4547-a3ec-5a58e234ce5e/resourceGroups/  
  demo-res-group/providers/Microsoft.Storage/storageAccounts/demostorage20170418
```

```
azurerm_storage_account.my_storage_account: Importing from ID ...  
azurerm_storage_account.my_storage_account: Import complete!  
  Imported azurerm_storage_account (ID: ...)  
azurerm_storage_account.my_storage_account: Refreshing state... (ID: ...)
```

Import success! The resources imported are shown above. These are now in your Terraform state. Import does not currently generate configuration, so you must do this next. If you do not create configuration for the above resources, then the next 'terraform plan' will mark them for destruction.

```
$ terraform state list  
azurerm_storage_account.my_storage_account
```

```
$ terraform state show azurerm_storage_account.my_storage_account  
id                = /subscriptions/e9b2ec19...  
account_kind      = Storage  
account_type      = Standard_LRS  
location          = westeurope  
name              = demostorage20170418  
...
```

Example: Use Remote State

```
$ terraform init -backend=true \  
-backend-config="backend=azure" \  
-backend-config="storage_account_name=demostorage20170418" \  
-backend-config="container_name=demo-storage-container" \  
-backend-config="key=network.terraform.tfstate"
```

Example: Use Remote State to Chain Projects

```
data "terraform_remote_state" "net" {  
  backend = "azure"  
  config {  
    storage_account_name = "demostorage20170418"  
    container_name       = "demo-storage-container"  
    key                  = "network.terraform.tfstate"  
  }  
}  
  
resource "azurerm_public_ip" "rhelvm" {  
  name          = "demo-rhelvm-ip"  
  location      =  
    "${data.terraform_remote_state.net.location}"  
  resource_group_name =  
    "${data.terraform_remote_state.net.resource_group_name}"  
  public_ip_address_allocation = "static"  
}
```

Example: Using Data Source to lookup data

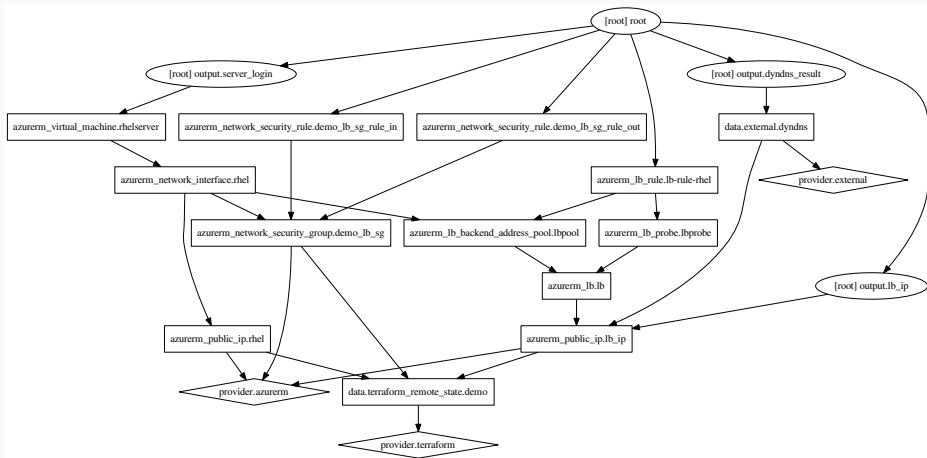
```
# searches for most recent tagged AMI in own account
data "aws_ami" "webami" {
  most_recent = true
  owners      = ["self"]

  filter {
    name     = "tag:my_key"
    values   = ["my_value"]
  }
}



# use AMI
resource "aws_instance" "web" {
  instance_type = "t2.micro"
  ami           = "${data.aws_ami.webami.id}"
}
```

Example: “External” Data Source

```
data "external" "dyndns" {  
  program = ["bash", "${path.module}/variomedia_dyndns.sh"]  
  
  query = {  
    hostname = "azure-demo.martin-schuetzte.de"  
    ipaddress = "${azurerm_public_ip.lb_ip.ip_address}"  
  }  
}
```



How to Write Own Plugins

- Learn you some Golang 
- Use the schema helper lib
- Adapt to model of
Provider (setup steps, authentication) and
Resources (arguments/attributes and CRUD methods)
- Start reading of simple plugins like
`builtin/providers/mysql` 

Usage

Under active development,
current version 0.9.3 (April 12), expecting 1.0 soon

- Modules are very simple
- Lacking syntactic sugar
(e.g. aggregations, common repetitions)
- Big improvements in state management
- Large variation in provider support

- Testing is inherently difficult
- Provider coverage largely depends on community
- Resource model mismatches, e.g. with Heroku apps
- Ignorant of API rate limits, account ressource limits, etc.

- Lists and Maps may be passed to modules
- State Import
- Data Sources
- State Environments

Configuration Management Tools:

- SaltStack Salt Cloud
- Ansible modules
- Puppet modules

Vendor Tools:

- Azure Resource Manager Templates
- AWS CloudFormation
- OpenStack Heat

- Avoid user credentials in Terraform code, use e.g. profiles and *assume-role* wrapper scripts
- At least use separate user credentials, know how to revoke them
- To hold credentials in VCS use PGP encryption, e.g. with [Blackbox](#)

- Use a VCS, i. e. git
- Always add some `"${var.shortname}"` to namespace
- Use remote state and consider access locking, e. g. with a single build server
- Take a look at [Hashicorp Atlas](#) and its workflow

PROVISION, SECURE, AND RUN

ANY INFRASTRUCTURE FOR ANY APPLICATION



LEARN THE HASHICORP SUITE >

PROVISION



Vagrant



Packer



Terraform

SECURE



Vault

RUN



Nomad



Consul

BUILD TEST

PACKAGE



PROVISION

SECURE

DEPLOY



MAINTAIN

Seven elements of the modern Application Lifecycle

- [Terraform.io](https://terraform.io) and [hashicorp/terraform](https://github.com/hashicorp/terraform) 
- [terraform-community-modules](https://github.com/terraform-community-modules) 
- [Creating Azure Resources with Terraform](#)
- [A Comprehensive Guide to Terraform](#)
- [Terraform, VPC, and why you want a tfstate file per env](#)
- [Terraform: Beyond the Basics with AWS](#)

Hopefully, deployments will become routine and boring—and in the world of operations, boring is a very good thing.

— *Terraform: Up & Running* by Yevgeniy Brikman



Defining system infrastructure as code and building it with tools doesn't make the quality any better. At worst, it can complicate things.

— *Infrastructure as Code* by Kief Morris



Martin Schütte
@m_schuett
info@martin-schuette.de

slideshare.net/mschuett/ 